



SCYLLA

The Real-Time Big Data Database

Introduction

Good [morning/afternoon] everyone. Today I'll be presenting an operational analysis of ScyllaDB, a high-performance NoSQL database system.

In this presentation, I'll cover four key areas of ScyllaDB's operational capabilities:

1. **CRUD Operations:** Demonstrating the fundamental Create, Read, Update, and Delete operations using ScyllaDB's CQL syntax
2. **Query Optimization Techniques:** Focusing on partition key selection and efficiency strategies
3. **Advanced Features and Transaction Management:** Including materialized views, batch operations, and time-to-live functionality

Performance Monitoring and Troubleshooting: Using built-in tools and external monitoring systems

SCYLLA DB

OPERATIONAL ANALYSIS

ScyllaDB Operational Analysis

```
[cqlsh> CREATE KEYSPACE tamu_aggies WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};

Warnings :
SimpleStrategy replication class is not recommended, but was used for keyspace tamu_aggies. You may suppress this warning by delisting SimpleStrategy from replication_strategy_warn_list config option or by disabling this replication strategy on replication_strategy_fail_list.

[cqlsh> use tamu_aggies;
cqlsh:tamu_aggies> CREATE TABLE my_class (
...     id UUID PRIMARY KEY,
...     name TEXT,
...     value INT
... );
[cqlsh:tamu_aggies>
[cqlsh:tamu_aggies>
cqlsh:tamu_aggies> CREATE TYPE address (
...     street TEXT,
...     city TEXT,
...     zip_code INT
... );
cqlsh:tamu_aggies> █
```

Part 1: CRUD Operations

Understanding Keyspaces

In ScyllaDB (and Cassandra), a keyspace is similar to a database schema in relational systems:

- Acts as a top-level namespace defining replication across nodes
- Contains tables, replication strategy, and replication factor

```
CREATE KEYSPACE family
WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
```

- **Replication Strategy:** Defines how data copies are distributed
- **Replication Factor:** Number of data copies (1 = development, 3+ = production)

Create Operations

```
CREATE KEYSPACE tamu_aggies
WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};
```

```
CREATE TABLE my_class (  
    id UUID PRIMARY KEY,  
    name TEXT,  
    value INT  
);
```

```
CREATE TYPE address (  
    street TEXT,  
    city TEXT,  
    zip_code INT  
);
```

Primary Keys in ScyllaDB

Primary keys:

- Must be unique for each row
- Cannot be null
- Determine data distribution across the cluster
- UUID is ideal as it's globally unique and can be generated client-side

Alter Operations

```
ALTER TABLE my_class ADD email TEXT;
```

```
ALTER TABLE my_class WITH compaction = {  
    'class': 'SizeTieredCompactionStrategy',  
    'max_threshold': 32  
};
```

```
cqlsh:tamu_aggies> ALTER TABLE my_class ADD email TEXT;  
cqlsh:tamu_aggies> ALTER TABLE my_class WITH compaction = {  
    ...    'class': 'SizeTieredCompactionStrategy',  
    ...    'max_threshold': 32  
    ... };
```

Drop Operations

```
USE tamu_aggies;  
DROP TABLE my_class;  
DROP KEYSPACE tamu_aggies;
```

```
DROP TYPE address;
```

```
[cqlsh> DROP KEYSPACE tamu_aggies;
[cqlsh> DROP TYPE address;
InvalidRequest: Error from server: code=2200 [Invalid query] message="No keyspace has been specified
. USE a keyspace, or explicitly specify keyspace.tablename"
[cqlsh> use tamu_aggies;
InvalidRequest: Error from server: code=2200 [Invalid query] message="Keyspace 'tamu_aggies' does not exist"
```

Secondary Indexes

Create additional access patterns beyond primary key lookups:

```
CREATE INDEX ON my_table(email);
CREATE INDEX my_custom_index ON my_table(age);

-- Using indexes
SELECT * FROM my_table WHERE email = 'user@example.com';
```

```
[cqlsh> CREATE KEYSPACE IF NOT EXISTS tamu_aggies
[ ... WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 3};

Warnings :
SimpleStrategy replication class is not recommended, but was used for keyspace tamu_aggies. You may suppress this warning by delisting SimpleStrategy from replication_strategy_warn_list configuration option, or make it into an error by listing this replication strategy on replication_strategy_fail_list.
```

```
cqlsh> use tamu_aggies;
cqlsh:tamu_aggies> CREATE TABLE course_enrollment (
[ ...
    course_id uuid PRIMARY KEY,
    course_name text,
    department text,
    max_capacity int,
    current_enrollment int
    ... );
cqlsh:tamu_aggies> INSERT INTO course_enrollment (course_id, course_name, department, max_capacity, current_enrollment)
[ ... VALUES (uuid(), 'Advanced Databases', 'ISTM', 70, 55);
cqlsh:tamu_aggies> INSERT INTO course_enrollment (course_id, course_name, department, max_capacity, current_enrollment)
    ... VALUES (uuid(), 'Data Mining', 'ISTM', 70, 55);
cqlsh:tamu_aggies> SELECT * FROM course_enrollment;
```

course_id	course_name	current_enrollment	department	max_capacity
f3204eb3-7f18-41fa-9c86-2d13976d5804	Data Mining	55	ISTM	70
32eee364-09cf-4c30-9a44-b8314a255ba8	Advanced Databases	55	ISTM	70

Part 2: Query Optimization

Partition Key Optimization

- Choose partition keys that evenly distribute data
- Well-designed keys improve read/write performance

- Example: UUID as partition key ensures even distribution

Lightweight Transactions

```
UPDATE course_enrollment
SET enrollment_count = enrollment_count - 1
WHERE course_id = [specific-uuid]
IF enrollment_count > 0;
```

```
cqlsh:tamu_aggies> UPDATE course_enrollment
... SET current_enrollment = 56
... WHERE course_id = 32eee364-09cf-4c30-9a44-b8314a255ba8
[
... IF current_enrollment = 55;
]
```

[applied]	current_enrollment
True	55

- Uses IF clause for compare-and-set operations
- Maintains data integrity in distributed environment
- Performance impact: milliseconds with partition key vs. seconds without

Part 3: Transaction Management

Batch Operations

```
cqlsh:tamu_aggies> BEGIN BATCH
... -- Update enrollment in Advanced Databases - using direct assignment instead
... UPDATE course_enrollment
... SET current_enrollment = 54 -- Assuming current value is 55, set it to 54
... WHERE course_id = 32eee364-09cf-4c30-9a44-b8314a255ba8;
...
... -- Insert a related course with your specified values
... INSERT INTO course_enrollment
... (course_id, course_name, department, max_capacity, current_enrollment)
... VALUES (uuid(), 'Advanced Database Topics', 'ISTM', 70, 50);
[
... APPLY BATCH;
```

Example batch demonstrating atomic operations:

```
BEGIN BATCH
UPDATE courses SET enrollment_count = enrollment_count - 1
WHERE course_id = [Advanced-DB-UUID];

INSERT INTO courses (course_id, name, max_capacity, enrollment_count)
VALUES (uuid(), 'New Related Course', 70, 50);
APPLY BATCH;
```


- Operations either complete fully or fail entirely
- Maintains data integrity across multiple operations

Part 4: Performance Monitoring

Built-in Tracing Tools

```
TRACING ON;  
SELECT * FROM course_enrollment WHERE course_id = [specific-uuid];
```

cqlsh:tamu_aggies> TRACING ON;
Now Tracing is enabled
cqlsh:tamu_aggies> SELECT * FROM course_enrollment WHERE course_id = 32eee364-09cf-4c30-9a44-b8314a255ba8;

course_id	course_name	current_enrollment	department	max_capacity
32eee364-09cf-4c30-9a44-b8314a255ba8	Advanced Databases	54	ISTM	70

(1 rows)

Tracing session: 868f8590-0c5d-11f0-9442-5ea1199cef4f

activity	source	source_elapsed	client	timestamp	
				Execute CQL3 query	2025-03-29 05:20:27.7530
00	172.17.0.2	0	172.17.0.2	Parsing a statement [shard 15]	2025-03-29 05:20:27.7534
09	172.17.0.2	1	172.17.0.2	Processing a statement for authenticated user: anonymous [shard 15]	2025-03-29 05:20:27.7534
88	172.17.0.2	80	172.17.0.2	Executing read query (reversed false) [shard 15]	2025-03-29 05:20:27.7535
86	172.17.0.2	177	172.17.0.2	Creating read executor for token -764619882566839294 with all: [172.17.0.2] targets: [172.17.0.2] repair decision: NONE [shard 15]	2025-03-29 05:20:27.7536
94	172.17.0.2	286	172.17.0.2	Creating never_speculating_read_executor - speculative retry is disabled or there are no extra replicas to speculate with [shard 15]	2025-03-29 05:20:27.7536
95	172.17.0.2	287	172.17.0.2	read_data: querying locally [shard 15]	2025-03-29 05:20:27.7536
99	172.17.0.2	291	172.17.0.2	Start querying singular range {{-764619882566839294, pk{001032eee36409cf4c309a44b8314a255ba8}}} [shard 3]	2025-03-29 05:20:27.7537
91	172.17.0.2	8	172.17.0.2	[reader concurrency semaphore user] admitted immediately [shard 3]	2025-03-29 05:20:27.7537
96	172.17.0.2	12	172.17.0.2	[reader concurrency semaphore user] executing read [shard 3]	2025-03-29 05:20:27.7537
99	172.17.0.2	15	172.17.0.2	Querying cache for range {{-764619882566839294, pk{001032eee36409cf4c309a44b8314a255ba8}}} and slice [[-inf, +inf]] [shard 3]	2025-03-29 05:20:27.7544
45	172.17.0.2	661	172.17.0.2	Page stats: 1 partition(s), 0 static row(s) (0 live, 0 dead), 1 clustering row(s) (1 live, 0 dead), 0 range tombstone(s) and 4 cell(s) (4 live, 0 dead) [shard 3]	2025-03-29 05:20:27.7544
75	172.17.0.2	692	172.17.0.2	Querying is done [shard 3]	2025-03-29 05:20:27.7544
86	172.17.0.2	703	172.17.0.2	Done processing - preparing a result [shard 15]	2025-03-29 05:20:27.7571
18	172.17.0.2	3710	172.17.0.2	Request complete	2025-03-29 05:20:27.7567
27	172.17.0.2	3727	172.17.0.2		

Tracing reveals:

- Total query execution time
- Internal steps and timings
- Node involvement
- Potential bottlenecks

```
SELECT * FROM SYSTEM_SIZE_ESTIMATES
```

```
cqlsh:tamu_aggies> SELECT * FROM system.size_estimates
... WHERE keyspace_name = 'tamu_aggies';
```

keyspace_name	table_name	range_start	range_end	mean_partition_size	partitions_count
---------------	------------	-------------	-----------	---------------------	------------------

```
cqlsh:tamu_aggies> SELECT * FROM system_traces.sessions;
```

session_id	client	command	coordinator	duration	parameters	request	request_size	response_size	started_at
username									
868f8590-0c5d-11f0-9442-5ea1199cef4f	172.17.0.2	QUERY	172.17.0.2	3727	{'consistency_level': 'ONE', 'page_size': '100', 'query': 'SELECT * FROM course_enrollment WHERE course_id = 32eee364-09cf-4c30-9a44-b8314a255ba8;', 'serial_consistency_level': 'SERIAL', 'user_timestamp': '1743225627752925'}	Execute CQL3 query	108	338	2025-03-29 05:20:27.753000+0000
anonymous									

(1 rows)

```
cqlsh:tamu_aggies> SELECT * FROM system_traces.events;
```

session_id	event_id	scylla_parent_id	scylla_span_id	activity	source	source_elapsed	thread
868f8590-0c5d-11f0-9442-5ea1199cef4f	868f958b-0c5d-11f0-9442-5ea1199cef4f	0	4415110950987	172.17.0.2		1	shard 15
868f8590-0c5d-11f0-9442-5ea1199cef4f	868f989f-0c5d-11f0-9442-5ea1199cef4f	0	4415110950987	172.17.0.2		80	shard 15
868f8590-0c5d-11f0-9442-5ea1199cef4f	868f9c70-0c5d-11f0-9442-5ea1199cef4f	0	4415110950987	172.17.0.2		177	shard 15
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fa0ae-0c5d-11f0-9442-5ea1199cef4f	0	4415110950987	172.17.0.2		286	shard 15
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fa0bd-0c5d-11f0-9442-5ea1199cef4f	0	4415110950987	172.17.0.2		287	shard 15
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fa0dc-0c5d-11f0-9442-5ea1199cef4f	0	4415110950987	172.17.0.2		291	shard 15
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fa477-0c5d-11f0-a357-5ea6199cef4f	4415110950987	124029003917579	172.17.0.2		8	shard 3
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fa4a5-0c5d-11f0-a357-5ea6199cef4f	4415110950987	124029003917579	172.17.0.2		12	shard 3
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fa4c4-0c5d-11f0-a357-5ea6199cef4f	4415110950987	124029003917579	172.17.0.2		15	shard 3
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fb02-0c5d-11f0-a357-5ea6199cef4f	4415110950987	124029003917579	172.17.0.2		661	shard 3
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fb2f-0c5d-11f0-a357-5ea6199cef4f	4415110950987	124029003917579	172.17.0.2		692	shard 3
868f8590-0c5d-11f0-9442-5ea1199cef4f	868fbf9c-0c5d-11f0-a357-5ea6199cef4f	4415110950987	124029003917579	172.17.0.2		703	shard 3
868f8590-0c5d-11f0-9442-5ea1199cef4f	8690266a-0c5d-11f0-9442-5ea1199cef4f	0	4415110950987	172.17.0.2		3710	shard 15

(13 rows)

```
cqlsh:tamu_aggies> TRACING OFF;
Disabled Tracing.
```

tamu_aggies	course_enrollment	-3649185051365608035	-3648679467481425086	0	0
tamu_aggies	course_enrollment	-371867932353059173	-281048287044735901	0	0
tamu_aggies	course_enrollment	-3783549675032091211	-3649185051365608035	0	0
tamu_aggies	course_enrollment	-3982219112397168481	-3783549675032091211	0	0
tamu_aggies	course_enrollment	-4006056807664523468	-3982219112397168481	0	0
tamu_aggies	course_enrollment	-4039919941292295443	-4006056807664523468	0	0
tamu_aggies	course_enrollment	-4143264996811021494	-4039919941292295443	0	0
tamu_aggies	course_enrollment	-4200676460757592165	-4143264996811021494	0	0
tamu_aggies	course_enrollment	-4385118449920616163	-4200676460757592165	0	0
tamu_aggies	course_enrollment	-4670969658236431222	-4385118449920616163	0	0

SELECT * FROM SYSTEM_TRACE_SESSIONS

External Monitoring

I was unable to perform the node tools or any type of monitoring because of run time errors and advising me that I did not have enough memory resources in the version of Scylla DB to perform the operation.

- Use `nodetool status` for cluster health
- Monitor for:
 - High latencies in trace results
 - Large partition sizes causing distribution problems
 - Resource constraints (CPU, memory, disk I/O)

Conclusion

ScyllaDB represents a powerful evolution in NoSQL database technology:

- Combines CQL syntax accessibility with exceptional distributed performance
- Balances performance with data integrity through proper query optimization
- Provides comprehensive monitoring for visibility into system performance

These features make ScyllaDB ideal for organizations requiring:

- High-throughput, low-latency operations at scale
- Performance, reliability, and observability for modern enterprise applications

Thank you for your time.

References

"Cqlsh: The CQL Shell." *ScyllaDB Open Source Documentation*, ScyllaDB, <https://opensource.docs.scylladb.com/stable/cql/cqlsh.html>. Accessed 30 Mar. 2025.

"ScyllaDB - Part 1 - An Introduction." *YouTube*, uploaded by Programming with Mosh, 17 Oct. 2023, https://youtu.be/MWFWM_LcouY?si=TqnqvmMRM2m0gGmu.

"ScyllaDB - Part 2 - Data Modeling Basics." *YouTube*, uploaded by Programming with Mosh, 18 Oct. 2023, <https://youtu.be/D3j5rSUSaCE?si=UtkYH27irhuZkA-o>.

"ScyllaDB - Part 3 - Collections, Counters and Frozen Types." *YouTube*, uploaded by Programming with Mosh, 19 Oct. 2023, <https://youtu.be/JPkrdWMVpPk?si=xNnx2n8K-BwR4aJC>.

"ScyllaDB - Part 4 - Indexing and Materialized Views." *YouTube*, uploaded by Programming with Mosh, 20 Oct. 2023, <https://youtu.be/cOlyquVuFJU?si=s0oidMDDiJZTnIJ3>.

"ScyllaDB - Part 5 - Lightweight Transactions." *YouTube*, uploaded by Programming with Mosh, 23 Oct. 2023, https://youtu.be/A5TkqoMxoJk?si=3zh75V_kUNEsCysI.